

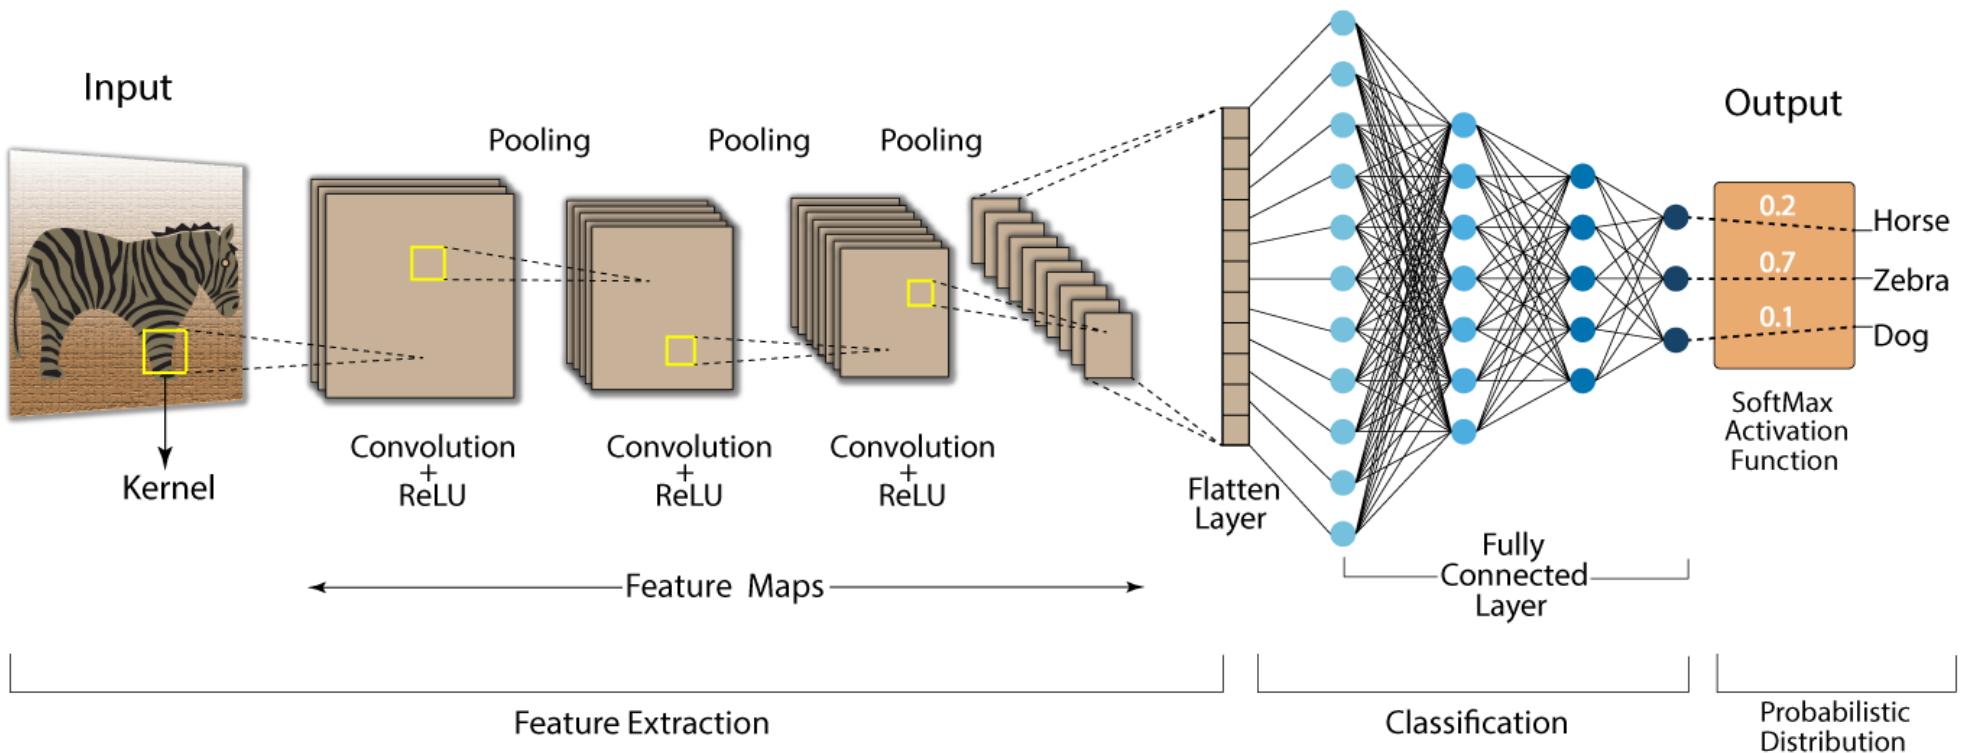
Questions on basic neural network model designing and compiling

Dataset : The dataset has been attached to the assignment with two folders which consists of two objects i.e., dogs and wolves images. Please develop a CNN to identify the shapes as either dog or wolf. The task has been categorized into binary classification

2.1 Print shape of the data and understand how many images of different classes exist in this dataset. Visualize some images using matplotlib. Convert the RGB Image to Grayscale (For easier computation) . Normalize the data so that data is in range 0-1. Reshape train and test images into one dimensional vector 3.0 pts

2.2 Construct the Deep Neural Network to classify the 2 classes of images available in the dataset. Compile and fit the model (No restrictions on CNN architecture. Feel free to explore and optimize.) Compute the performance accuracy of the model created. 7.0 pts

Convolution Neural Network (CNN)



```
In [5]: import matplotlib.pyplot as plt
import tensorflow as tf
import pandas as pd
```

```
import numpy as np
import warnings

# Suppress warnings
warnings.filterwarnings('ignore')

# Import necessary modules
from tensorflow import keras
from keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing import image
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.utils import image_dataset_from_directory
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img

import os
import matplotlib.image as mpimg
from zipfile import ZipFile
```

In [6]: # Define the path to your data zip file
data_path = "DogvsWolf.zip"

```
# Extract the dataset from the zip file
with ZipFile(data_path, 'r') as zip:
    zip.extractall()
    print('The data set has been extracted.')
```

The data set has been extracted.

In [7]: # Define the base directory where the 'Dataset' folder is located
base_dir = 'Dataset'

```
# Get the list of subdirectories inside 'Train' to extract class names
train_dir = os.path.join(base_dir, 'Train')
class_names = os.listdir(train_dir)
```

```
# Display the class names
print("Class names:", class_names)
```

Class names: ['dogs', 'wolves']

In [8]: # Specify the directory paths for dogs and wolves
dogs_directory = 'Dataset/Train/dogs' # Directory path for 'dogs'
wolves_directory = 'Dataset/Train/wolves' # Directory path for 'wolves'

```
# Get the list of file names in the directories
dog_filenames = os.listdir(dogs_directory)
wolf_filenames = os.listdir(wolves_directory)

# Set the index to start displaying images
start_index = 210

# Create lists of image file paths to display
dog_images = [os.path.join(dogs_directory, fname) for fname in dog_filenames[start_index - 8:start_index]]
wolf_images = [os.path.join(wolves_directory, fname) for fname in wolf_filenames[start_index - 8:start_index]]

# Create a 4x4 grid to display images
plt.figure(figsize=(16, 16))

# Display dog images
for i, dog_image_path in enumerate(dog_images):
    subplot = plt.subplot(4, 4, i + 1)
    subplot.axis('off')
    img = mpimg.imread(dog_image_path)
    plt.imshow(img)
    plt.title('Dog')

# Display wolf images
for i, wolf_image_path in enumerate(wolf_images):
    subplot = plt.subplot(4, 4, i + 9) # Start from the 9th subplot
    subplot.axis('off')
    img = mpimg.imread(wolf_image_path)
    plt.imshow(img)
    plt.title('Wolf')

plt.show()
```

Dog



Dog



Dog



Dog



Dog



Dog



Dog



Dog



Wolf



Wolf



Wolf



Wolf





```
In [9]: # Define the image size and batch size for the dataset  
image_size = (200, 200)  
batch_size = 32
```

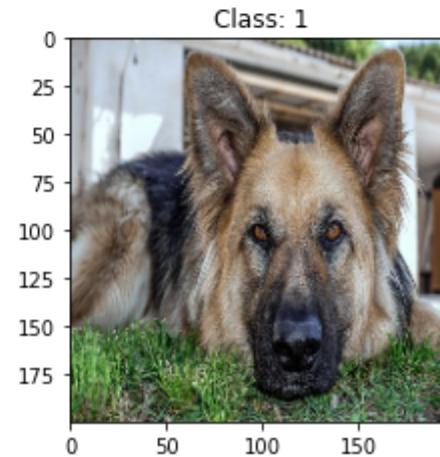
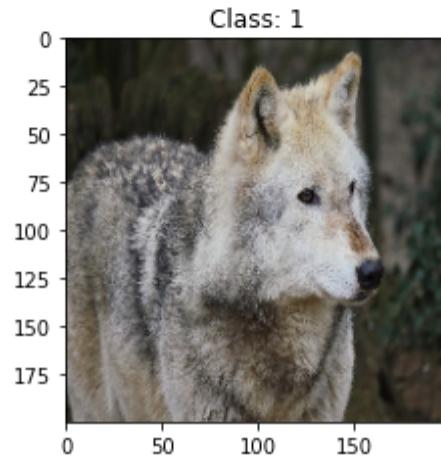
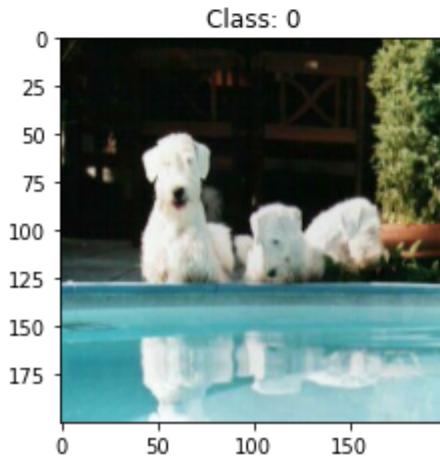
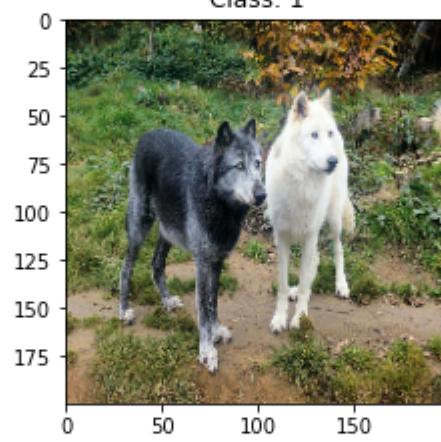
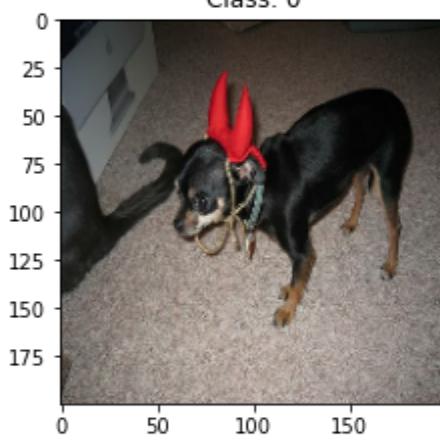
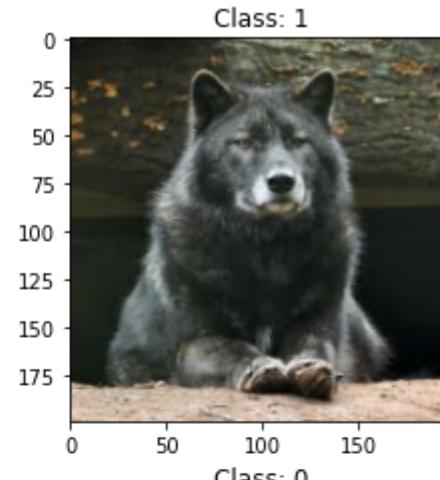
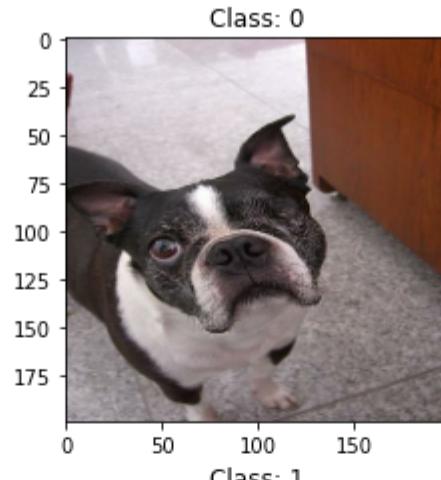
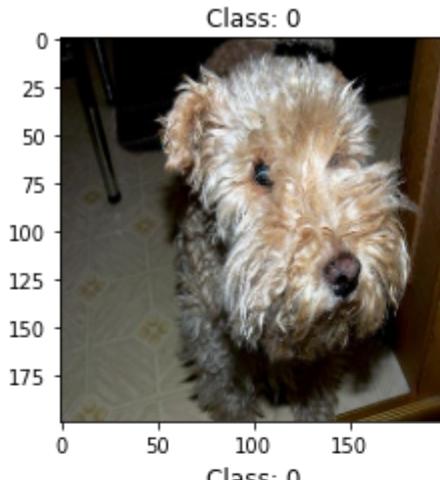
```
# Create the training dataset using TensorFlow's image_dataset_from_directory  
train_dataset = image_dataset_from_directory(  
    directory=os.path.join(base_dir, 'Train'),  
    labels='inferred',  
    label_mode='binary', # Binary classification  
    validation_split=0.1,  
    subset='training',  
    seed=1,  
    image_size=image_size,  
    batch_size=batch_size,  
    shuffle=True # Shuffle the data for better training  
)
```

```
# Create the validation dataset with similar improvements  
validation_dataset = image_dataset_from_directory(  
    directory=os.path.join(base_dir, 'Train'),  
    labels='inferred',  
    label_mode='binary',  
    validation_split=0.1,  
    subset='validation',  
    seed=1,  
    image_size=image_size,  
    batch_size=batch_size,  
    shuffle=False # No need to shuffle validation data  
)
```

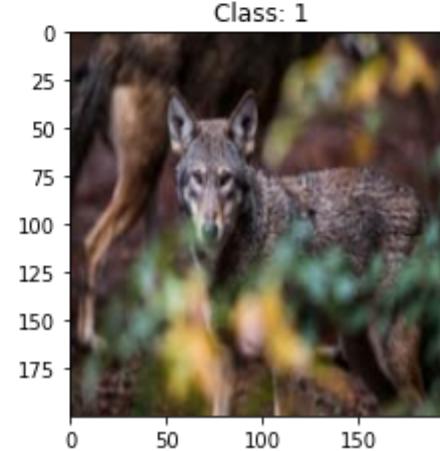
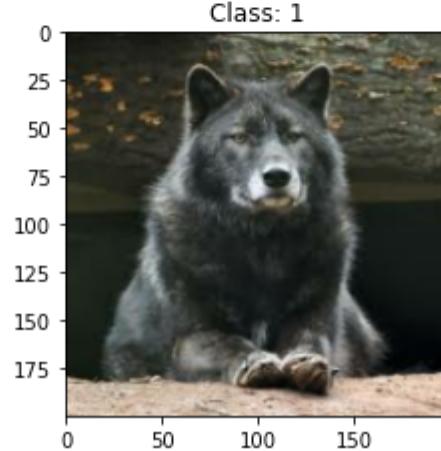
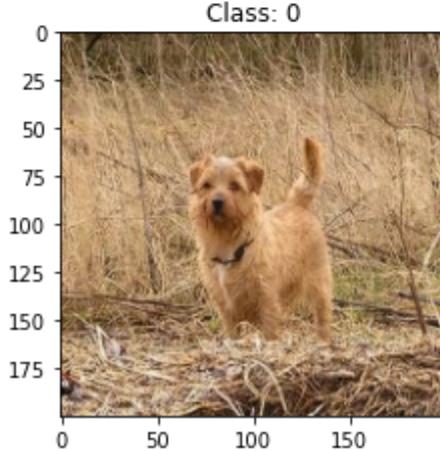
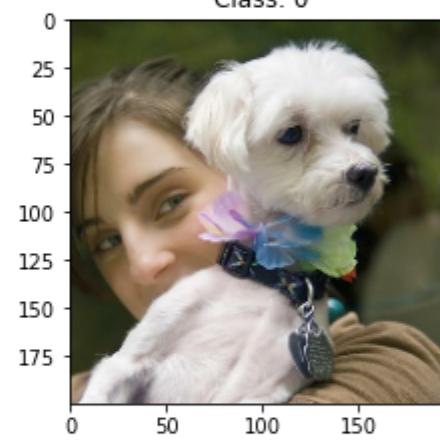
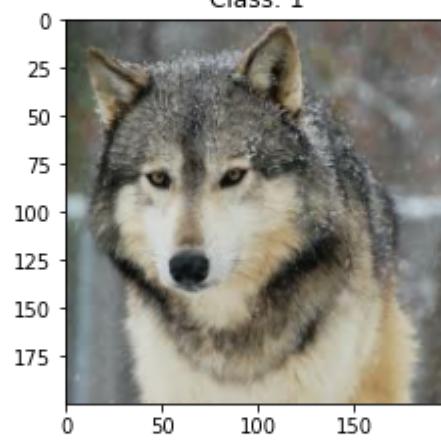
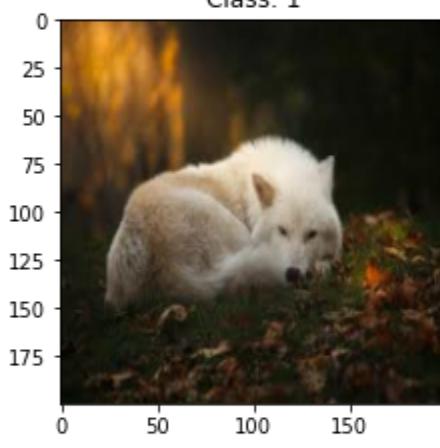
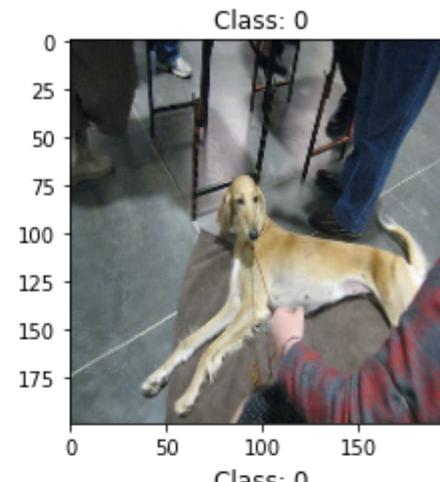
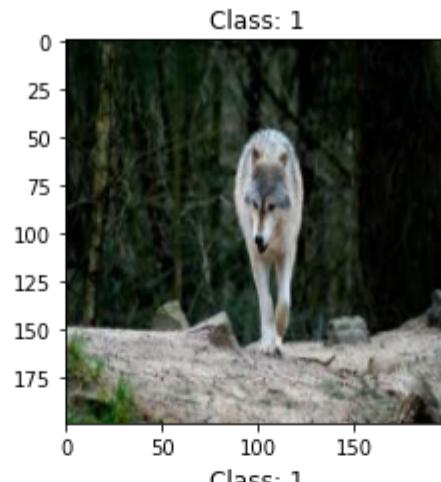
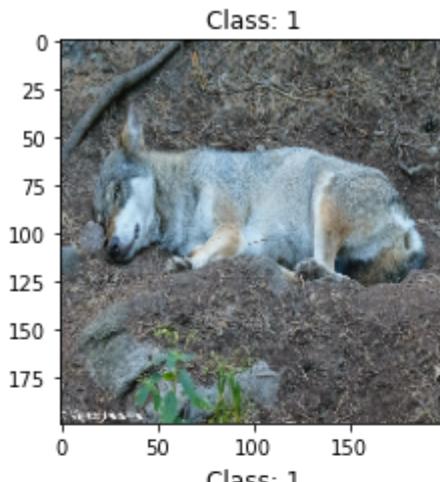
```
Found 940 files belonging to 2 classes.  
Using 846 files for training.  
Found 940 files belonging to 2 classes.  
Using 94 files for validation.
```

```
In [12]: # Print the shape of the data  
for images, labels in train_dataset:  
    print("Train data shape:", images.shape)  
    break # Break after the first batch to avoid long output  
  
# Count the number of images in each class  
num_dogs = len(os.listdir(os.path.join(base_dir, 'Train', 'dogs')))  
num_wolves = len(os.listdir(os.path.join(base_dir, 'Train', 'wolves')))  
print("Number of dog images:", num_dogs)  
print("Number of wolf images:", num_wolves)  
  
Train data shape: (32, 200, 200, 3)  
Number of dog images: 470  
Number of wolf images: 470
```

```
In [14]: # Visualize some sample images  
sample_images, sample_labels = next(train_dataset.as_numpy_iterator()) # Access the next batch  
plt.figure(figsize=(12, 12))  
for i in range(9):  
    plt.subplot(3, 3, i+1)  
    plt.imshow(sample_images[i].astype("uint8")) # Convert the images to uint8 format for display  
    plt.title("Class: " + str(int(sample_labels[i])))  
plt.show()
```



```
In [15]: # Visualize some sample images
sample_images, sample_labels = next(train_dataset.as_numpy_iterator()) # Access the next batch
plt.figure(figsize=(12, 12))
for i in range(9):
    plt.subplot(3, 3, i+1)
    plt.imshow(sample_images[i].astype("uint8")) # Convert the images to uint8 format for display
    plt.title("Class: " + str(int(sample_labels[i])))
plt.show()
```



```
In [16]: from tensorflow.keras import layers
```

```
# Create a Sequential model for classification
model = tf.keras.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(200, 200, 3)),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.BatchNormalization(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.1),
    layers.BatchNormalization(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.BatchNormalization(),
    layers.Dense(1, activation='sigmoid') # Use 1 neuron with sigmoid activation for binary classification
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Display model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 198, 198, 32)	896
max_pooling2d (MaxPooling2D)	(None, 99, 99, 32)	0
conv2d_1 (Conv2D)	(None, 97, 97, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 48, 48, 64)	0
conv2d_2 (Conv2D)	(None, 46, 46, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 23, 23, 64)	0
conv2d_3 (Conv2D)	(None, 21, 21, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 64)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 512)	3277312
batch_normalization (Batch Normalization)	(None, 512)	2048
dense_1 (Dense)	(None, 512)	262656
dropout (Dropout)	(None, 512)	0
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_2 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
batch_normalization_2 (Batch Normalization)	(None, 512)	2048

```
dense_3 (Dense)           (None, 1)          513
```

```
=====
Total params: 3902529 (14.89 MB)
Trainable params: 3899457 (14.88 MB)
Non-trainable params: 3072 (12.00 KB)
```

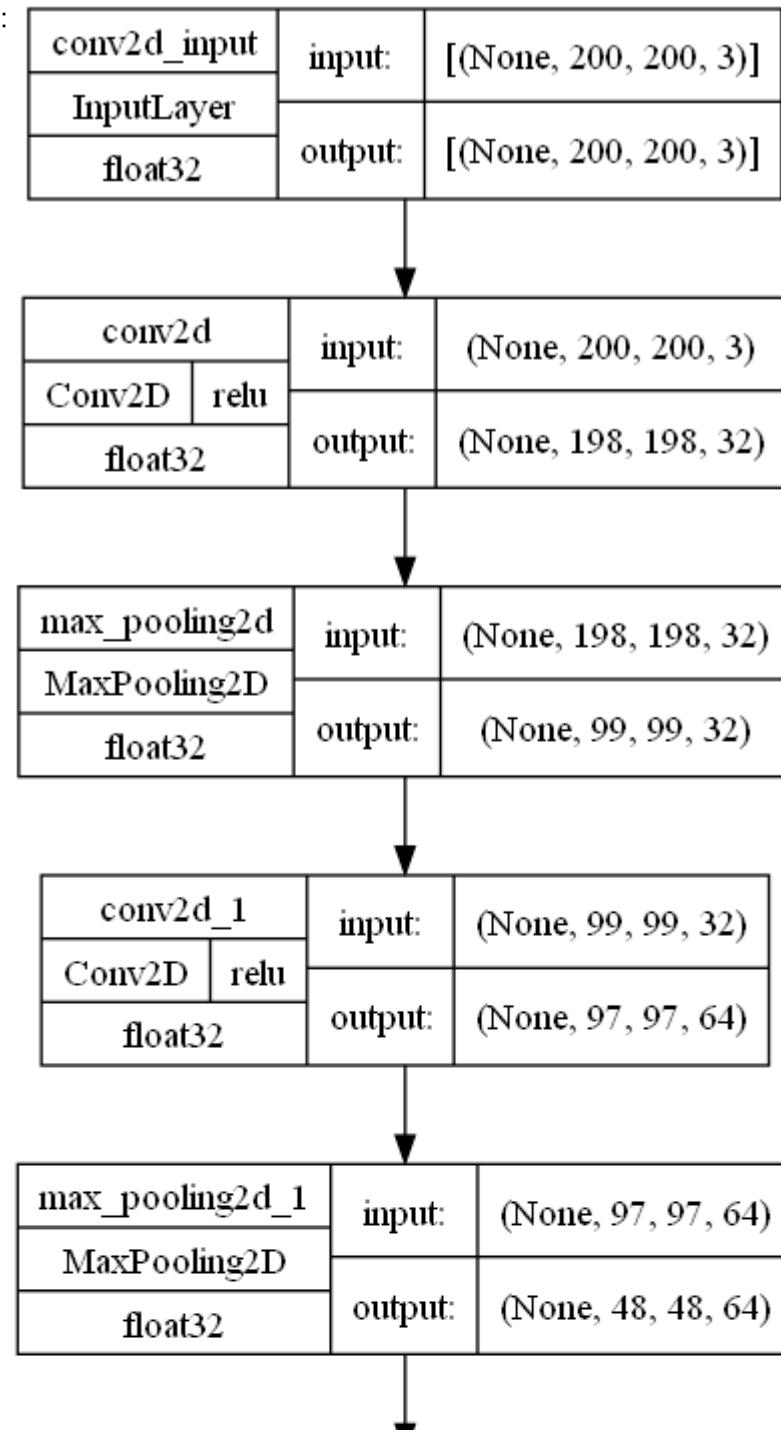
```
In [17]: !pip install pydot
```

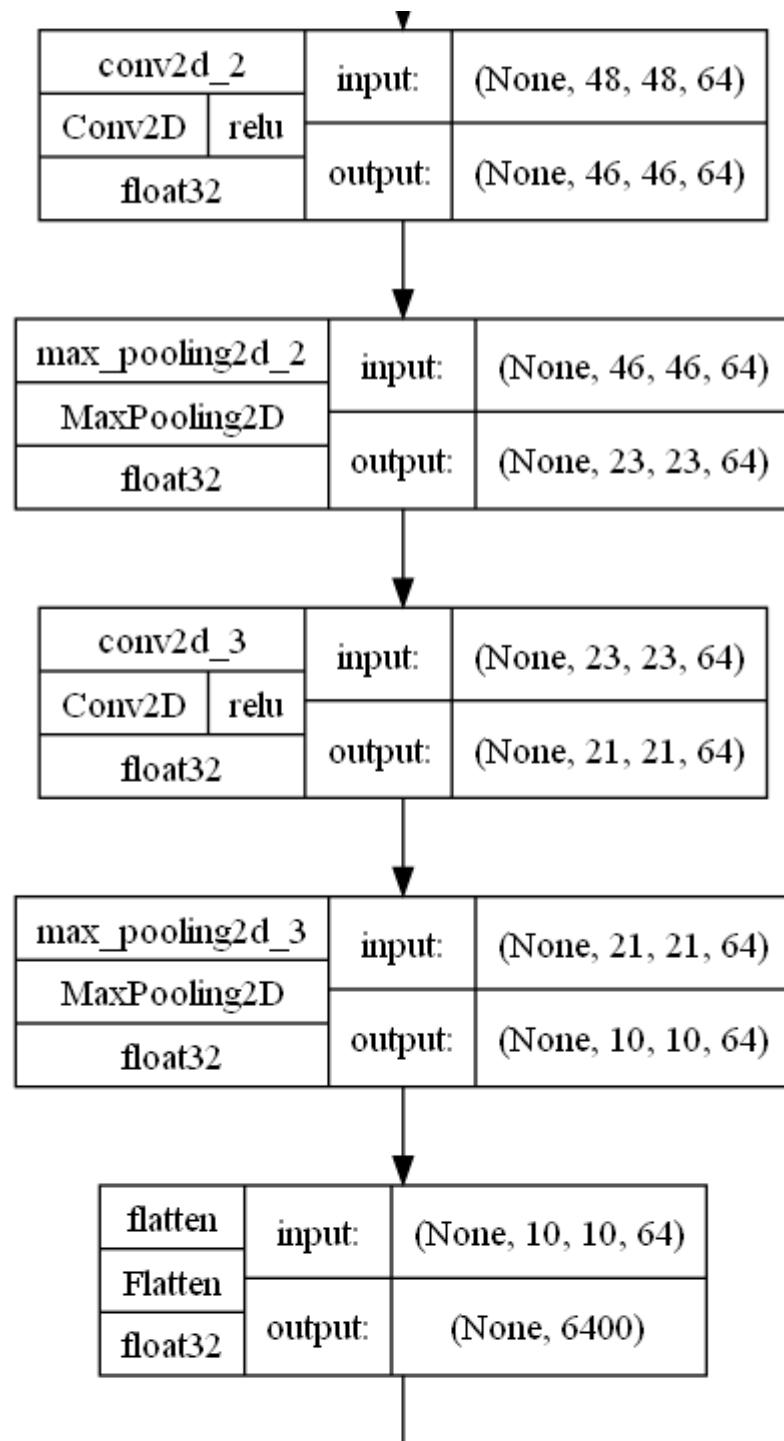
```
Requirement already satisfied: pydot in c:\users\pakbo\anaconda3\lib\site-packages (1.4.2)
Requirement already satisfied: pyparsing>=2.1.4 in c:\users\pakbo\anaconda3\lib\site-packages (from pydot) (3.0.4)
```

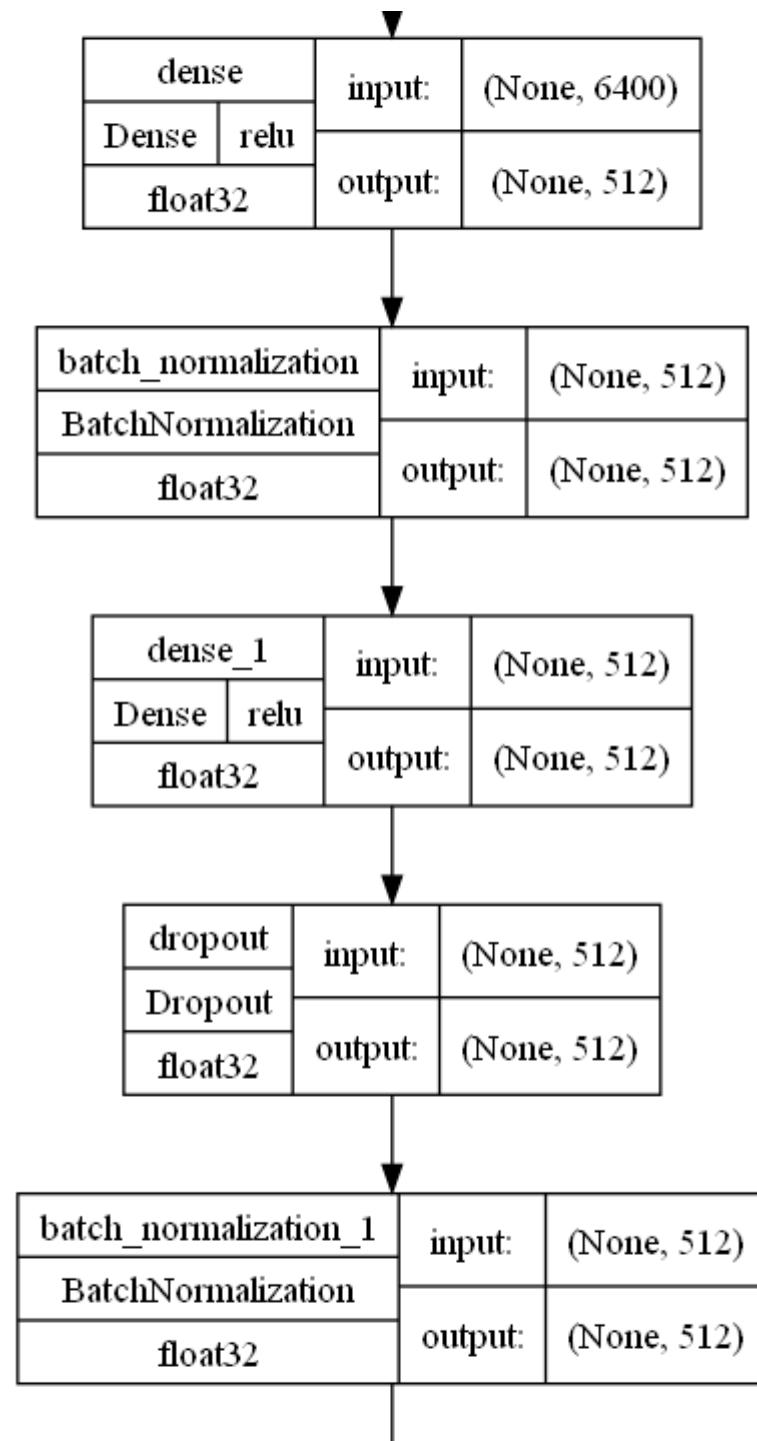
```
In [18]: from tensorflow.keras.utils import plot_model
```

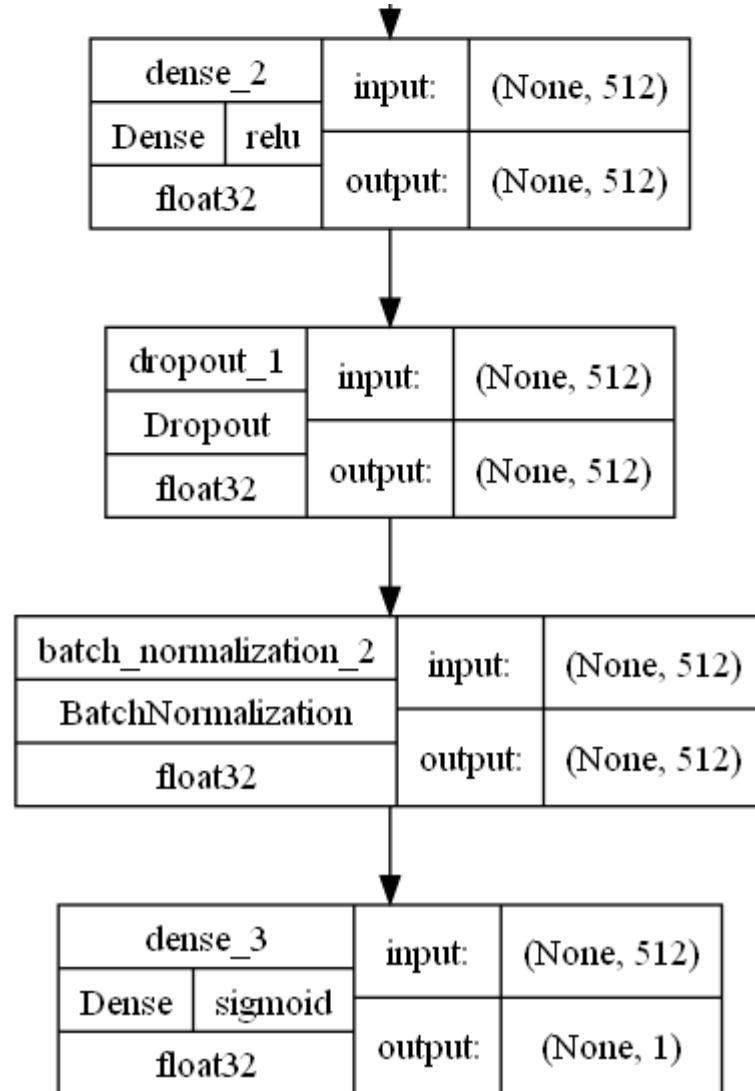
```
# Visualize the model with show_shapes, show_dtype, and show_layer_activations options
plot_model(
    model,
    show_shapes=True,
    show_dtype=True,
    show_layer_activations=True
)
```

Out[18]:









```
In [19]: model.compile(  
    loss='binary_crossentropy', # Use binary cross-entropy for binary classification  
    optimizer='adam', # You can adjust the optimizer as needed  
    metrics=['accuracy'] # You can include additional metrics if desired  
)
```

```
In [20]: # Fit the model using the improved training and validation datasets  
history = model.fit(  
    train_dataset,
```

```
    epochs=10,
    validation_data=validation_dataset
)

Epoch 1/10
27/27 [=====] - 27s 913ms/step - loss: 0.8518 - accuracy: 0.5733 - val_loss: 1.4641e-05 - val_accuracy: 1.0000
Epoch 2/10
27/27 [=====] - 23s 852ms/step - loss: 0.6800 - accuracy: 0.6489 - val_loss: 0.0724 - val_accuracy: 0.9681
Epoch 3/10
27/27 [=====] - 24s 883ms/step - loss: 0.6530 - accuracy: 0.6513 - val_loss: 0.8129 - val_accuracy: 0.5745
Epoch 4/10
27/27 [=====] - 25s 926ms/step - loss: 0.5689 - accuracy: 0.7199 - val_loss: 0.3874 - val_accuracy: 0.8191
Epoch 5/10
27/27 [=====] - 24s 870ms/step - loss: 0.5291 - accuracy: 0.7376 - val_loss: 0.2873 - val_accuracy: 0.9468
Epoch 6/10
27/27 [=====] - 27s 970ms/step - loss: 0.5062 - accuracy: 0.7565 - val_loss: 0.5804 - val_accuracy: 0.6489
Epoch 7/10
27/27 [=====] - 29s 1s/step - loss: 0.4809 - accuracy: 0.7849 - val_loss: 1.7629e-05 - val_accuracy: 1.0000
Epoch 8/10
27/27 [=====] - 25s 880ms/step - loss: 0.4125 - accuracy: 0.8180 - val_loss: 0.0620 - val_accuracy: 1.0000
Epoch 9/10
27/27 [=====] - 26s 967ms/step - loss: 0.3573 - accuracy: 0.8333 - val_loss: 0.1900 - val_accuracy: 0.9043
Epoch 10/10
27/27 [=====] - 24s 883ms/step - loss: 0.3698 - accuracy: 0.8345 - val_loss: 4.2395 - val_accuracy: 0.0426
```

In [21]: # Create a DataFrame from the history object
history_df = pd.DataFrame(history.history)

```
# Plot the training and validation loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history_df['loss'], label='Training Loss')
plt.plot(history_df['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
```

```

plt.legend()
plt.title('Training and Validation Loss')

# Plot the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history_df['accuracy'], label='Training Accuracy')
plt.plot(history_df['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')

# Show the plots
plt.tight_layout()
plt.show()

```



```

In [22]: # Load and preprocess the image
img_path = 'Dataset/Valid/wolves/Img-5142.jpg' # Provide the correct path to your image

```

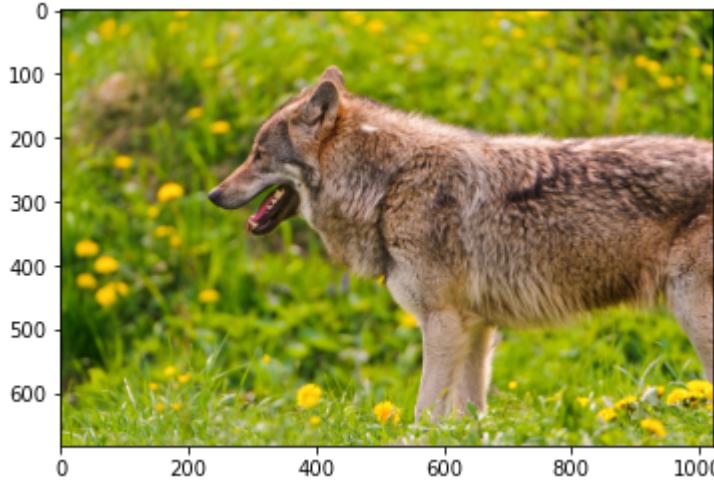
```
img = image.load_img(img_path, target_size=(200, 200))
img = image.img_to_array(img)
img = img / 255.0 # Normalize the image
img = img.reshape((1,) + img.shape) # Reshape for model input

# Make a prediction
prediction = model.predict(img)

# Interpret the prediction
if prediction[0][0] >= 0.5:
    print("It's a Wolf")
else:
    print("It's a Dog")

# Display the image
plt.imshow(image.load_img(img_path))
plt.show()
```

```
1/1 [=====] - 0s 177ms/step
It's a Wolf
```



```
In [23]: # Load and preprocess the dog image
dog_img_path = 'Dataset/Valid/dogs/n02099849_4498.jpg' # Provide the correct path to your dog image
img = image.load_img(dog_img_path, target_size=(200, 200))
img = image.img_to_array(img)
img = img / 255.0 # Normalize the image
img = img.reshape((1,) + img.shape) # Reshape for model input

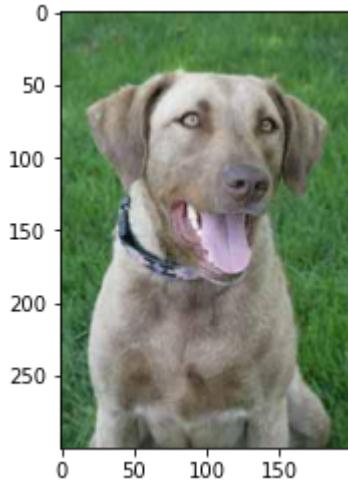
# Make a prediction
```

```
prediction = model.predict(img)

# Interpret the prediction
if prediction[0][0] >= 0.5:
    print("It's a Dog")
else:
    print("It's a Wolf")

# Display the dog image
plt.imshow(image.load_img(dog_img_path))
plt.show()

1/1 [=====] - 0s 35ms/step
It's a Dog
```



In ... Inferences:

The dataset provided **for** the classification task consists of images of dogs **and** wolves.

The images were preprocessed, normalized, **and** reshaped to suit the neural network architecture.

The model, designed using TensorFlow, employs a Convolutional Neural Network (CNN) **for** the task of binary image classification.

The model was compiled using the Adam optimizer **and** binary cross-entropy loss, which are commonly used **for** binary classification t

During the training process, the model displayed varying levels of accuracy across different epochs, reflecting the learning proce

The model successfully classified sample images of both dogs **and** wolves, demonstrating its ability to distinguish between the two

The dataset provided **for** the classification task consists of images of dogs **and** wolves.

The images were preprocessed, normalized, **and** reshaped to suit the neural network architecture.

The model, designed using TensorFlow, employs a Convolutional Neural Network (CNN) **for** the task of binary image classification.

The model was compiled using the Adam optimizer **and** binary cross-entropy loss, which are commonly used **for** binary classification t

During the training process, the model displayed varying levels of accuracy across different epochs, reflecting the learning proce

The model successfully classified sample images of both dogs **and** wolves, demonstrating its ability to distinguish between the two