DEEP LEARNING TENSOR OPS

Input Image — Convolution Layers — Pooling Layers — Fully Connected Layers — Output

Output classes: aeroplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck

Dog — Wolf

Mini Project By

Sangeetha Shankar

G.R.Priyavarshini

Shafeena Farheen

1.1 (i) Write a program to create a variable with a size of 2x3 and a constant (2 marks)

Create a TensorFlow variable of size 2x3 and a TensorFlow constant of the same size. The variables are then displayed using the "print" statements. The output shows the values of the TensorFlow variable and constant.

1.1 Tensor Operations:

(i) Write a program to create a variable with a size of 2x3 and a constant:

```python
import tensorflow as tf

# Create a TensorFlow variable with a size of 2x3
tf_variable = tf.Variable([[1, 2, 3], [4, 5, 6]])

# Create a TensorFlow constant with the same size
tf_constant = tf.constant([[1, 1, 1], [1, 1, 1]])

# Display the variables
print("TensorFlow Variable:")
print(tf_variable.numpy())
print("TensorFlow Constant:")
print(tf_constant.numpy())
```

```
TensorFlow Variable:
[[1 2 3]
 [4 5 6]]
TensorFlow Constant:
[[1 1 1]
 [1 1 1]]
```

Inference: The variable "tf_variable" is a 2x3 matrix with values [[1, 2, 3], [4, 5, 6]], while the constant "tf_constant" is a 2x3 matrix with values [[1, 1, 1], [1, 1, 1]].

(ii) Write a program to create a variable "tf3" with a size 3x3 and convert it to a numpy size of 3x3. Display rank of the matrix (2 marks)

The code in question 2 creates a TensorFlow variable "tf3" with a size of 3x3. It then converts the TensorFlow variable to a NumPy array using the "numpy()" function. The rank of the matrix is calculated using the "np.linalg.matrix_rank()" function and displayed using the "print" statement.

(ii) Write a program to create a variable "tf3" with a size 3x3 and convert it to a numpy size of 3x3. Display the rank of the matrix:

```python
import tensorflow as tf
import numpy as np

# Create a TensorFlow variable with a size of 3x3
tf3 = tf.Variable([[1, 2, 3], [4, 5, 6], [7, 8, 9]])

# Convert the TensorFlow variable to a NumPy array
numpy_array = tf3.numpy()

# Calculate and display the rank of the matrix
rank = np.linalg.matrix_rank(numpy_array)
print("NumPy Array:")
print(numpy_array)
print("Matrix Rank:", rank)
```

```
NumPy Array:
[[1 2 3]
 [4 5 6]
 [7 8 9]]
Matrix Rank: 2
```

Inference: The variable "tf3" is a 3x3 matrix with values [[1, 2, 3], [4, 5, 6], [7, 8, 9]]. The NumPy array representation of the matrix is [[1, 2, 3], [4, 5, 6], [7, 8, 9]]. The rank of the matrix is 2.

(iii) Write a program to randomly generate 3x3 matrix and assign 0 to the first row first column element (1 mark)

The code in question 3 generates a random 3x3 matrix using the "tf.random.uniform()" function. It then assigns the value 0 to the first row's first column element using the "tf.tensor_scatter_nd_update()" function. The modified matrix is converted to a NumPy array using the "numpy()" function and displayed using the "print" statement.

(iii) Write a program to randomly generate a 3x3 matrix and assign 0 to the first row's first column element:

```python
import tensorflow as tf
import numpy as np

# Randomly generate a 3x3 matrix
random_matrix = tf.random.uniform(shape=(3, 3))

# Assign 0 to the first row's first column element
random_matrix = tf.tensor_scatter_nd_update(random_matrix, indices=[[0, 0]], updates=[0.0])

# Convert the TensorFlow tensor to a NumPy array
numpy_array = random_matrix.numpy()

# Display the modified matrix
print("Modified NumPy Array:")
print(numpy_array)
```

```
Modified NumPy Array:
[[0.         0.49558318 0.76333857]
 [0.13806903 0.2510687  0.5553553 ]
 [0.46918166 0.84995985 0.70758307]]
```

Inference: The modified NumPy array has the first element in the first row (index [0, 0]) as 0, while the rest of the elements are randomly generated.

1.2 (i). What is Cross Entropy? (2 marks)

1.2 Questions on Tensor Operations:

(i) What is Cross Entropy?

Definition: Cross entropy is a mathematical measure used to quantify the dissimilarity or difference between two probability distributions. It is commonly applied in various fields, including machine learning and statistics.

Formula: The formula for cross entropy, denoted as H(p, q), is expressed as:

$H(p, q) = -\sum(p(x) * log(q(x)))$

Where:

H(p, q) represents the cross entropy between two distributions.

p(x) is the probability of an event in the actual (true) distribution.

q(x) is the probability of the same event in the predicted distribution.

Cross entropy is a mathematical measure used to quantify the dissimilarity or difference between two probability distributions. It is commonly applied in various fields, including machine learning and statistics. The formula for cross entropy, denoted as H(p, q), is given by $H(p, q) = -\sum(p(x) * log(q(x)))$, where H(p, q) represents the cross entropy between two distributions, p(x) is the probability of an event in the actual (true) distribution, and q(x) is the probability of the same event in the predicted distribution.

Binary Cross Entropy (log loss): $H(p, q) = - [p * log(q) + (1 - p) * log(1 - q)]$

Categorical Cross Entropy: $H(p, q) = - \Sigma (p\_i * log(q\_i))$

In these formulas:

- p_i represents the true probability of the i-th class (in the case of categorical cross entropy).
- q_i represents the predicted probability of the i-th class (the model's output probability for that class).
- The summation is performed over all classes.

The goal during training is to minimize the cross-entropy loss. When p and q are identical, the cross entropy is minimized and equals zero. As the predicted probabilities (q) diverge from the true probabilities (p), the cross-entropy increases, indicating a larger error in the model's predictions. This makes cross entropy a suitable loss function for training classification models because it penalizes both overconfidence (when the model is too certain about incorrect predictions) and under confidence (when the model is uncertain about correct predictions).

-

(ii). Write a program to create a variable "tensor" with a size 5x8 and reshape it to a size of 4x10 (3 marks)

The code in question 5 creates a TensorFlow variable "tensor" with a size of 5x8 using the "tf.random.normal()" function. It then reshapes the tensor to a size of 4x10 using the "tf.reshape()" function. The reshaped tensor is displayed using the "print" statement.

(ii) Write a program to create a variable "tensor" with a size 5x8 and reshape it to a size of 4x10:

```
import tensorflow as tf

# Create a TensorFlow variable with a size of 5x8
tensor = tf.Variable(tf.random.normal((5, 8)))

# Reshape the tensor to a size of 4x10
reshaped_tensor = tf.reshape(tensor, (4, 10))

# Display the reshaped tensor
print("Reshaped Tensor:")
print(reshaped_tensor.numpy())
```

```
Reshaped Tensor:
[[ 0.6176739   0.17967848 -0.6753632   0.82978255 -0.85082877 -0.15073903
  -0.6687701   0.71692085  1.0233076  -1.0298293 ]
 [ 0.432221    0.3004995  -2.3387713   0.30415332 -0.27854866  1.0275682
  -0.45517752 -0.73647887 -0.58317876 -0.35651022]
 [-0.30280757  0.20946066 -0.47114554 -1.2850192   0.30444983  0.37857938
   1.0514244   1.34697    -0.8892283   0.65124536]
 [-0.34211546 -1.0141863   0.5970262  -0.92258257  0.25507808  0.90264654
   0.21576567 -0.25510165  0.6259984   1.3449931 ]]
```

Inference: The initial variable "tensor" is a 5x8 matrix with randomly generated values. After reshaping, the tensor becomes a 4x10 matrix with the same values but a different structure.

-

2 Questions on basic neural network model designing and compiling

Dataset:

The dataset has been attached to the assignment with two folders which consists of two objects i.e., dogs and wolves images. Please develop a CNN to identify the shapes as either dog or wolf. The task has been categorized into binary classification

Dataset for model building: Dataset with dogs and wolves images

```python
import matplotlib.pyplot as plt
import tensorflow as tf
import pandas as pd
import numpy as np
import warnings

# Suppress warnings
warnings.filterwarnings('ignore')

# Import necessary modules
from tensorflow import keras
from keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing import image
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.utils import image_dataset_from_directory
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img

import os
import matplotlib.image as mpimg
from zipfile import ZipFile
```

Import Libraries:

- **The code sets up the environment for deep learning and imports necessary libraries like TensorFlow, Keras, NumPy, and Pandas.**

- **It suppresses warning messages during execution to ensure a cleaner output.**

- **It involves handling and preprocessing image data, likely for building an image classification model.**

- **The code imports modules for image data augmentation, dataset management, and visualization using matplotlib.**

- **It includes functions for handling zip files, potentially for dataset extraction or storage.**

- **The use of TensorFlow and Keras functions indicates the creation of a convolutional neural network (CNN) for image classification.**

- **Data generators and the 'image_dataset_from_directory' function are used to prepare image datasets for model training.**

-

- **The code structure suggests the development and training of a CNN for image classification tasks, possibly for a specific application or project.**

```python
# Define the path to your data zip file
data_path = "DogvsWolf.zip"

# Extract the dataset from the zip file
with ZipFile(data_path, 'r') as zip:
    zip.extractall()
    print('The data set has been extracted.')

The data set has been extracted.
```

Data Extraction:

- The code defines the path to the data zip file and uses the ZipFile module to extract the dataset.

- The extracted dataset is stored in the "Dataset" folder.

- The code successfully extracts the dataset from the provided zip file and saves it in the "Dataset" folder.

```python
# Define the base directory where the 'Dataset' folder is located
base_dir = 'Dataset'

# Get the list of subdirectories inside 'Train' to extract class names
train_dir = os.path.join(base_dir, 'Train')
class_names = os.listdir(train_dir)

# Display the class names
print("Class names:", class_names)

Class names: ['dogs', 'wolves']
```

Class Names:

- The code retrieves the list of subdirectories inside the "Train" directory.

- The names of these subdirectories are considered as the class names for classification.

- The output displays the names of the classes present in the dataset.

- In this case, the classes are "dogs" and "wolves".

-
- - **The code retrieves the subdirectories inside the "Train" directory.**

- - **The subdirectory names are considered as the class names for classification.**

2.1 Print shape of the data and understand how many images of different classes exist in this dataset. Visualize some images using matplotlib. Convert the RGB Image to Grayscale (For easier computation). Normalize the data so that data is in range 0-1. Reshape train and test images into one dimensional vector   3.0 pts

```python
# Specify the directory paths for dogs and wolves
dogs_directory = 'Dataset/Train/dogs'  # Directory path for 'dogs'
wolves_directory = 'Dataset/Train/wolves'  # Directory path for 'wolves'

# Get the list of file names in the directories
dog_filenames = os.listdir(dogs_directory)
wolf_filenames = os.listdir(wolves_directory)

# Check the number of images in each class
num_dogs_train = len(os.listdir(dogs_directory))
num_wolves_train = len(os.listdir(wolves_directory))
print("Number of images in the 'dogs' class in the training dataset:", num_dogs_train)
print("Number of images in the 'wolves' class in the training dataset:", num_wolves_train)

# Set the index to start displaying images
start_index = 210

# Create lists of image file paths to display
dog_images = [os.path.join(dogs_directory, fname) for fname in dog_filenames[start_index - 8:start_index]]
wolf_images = [os.path.join(wolves_directory, fname) for fname in wolf_filenames[start_index - 8:start_index]]

# Create a 4x4 grid to display images
plt.figure(figsize=(16, 16))

# Display dog images
for i, dog_image_path in enumerate(dog_images):
    subplot = plt.subplot(4, 4, i + 1)
    subplot.axis('off')
    img = mpimg.imread(dog_image_path)
    plt.imshow(img)
    plt.title('Dog')

# Display wolf images
for i, wolf_image_path in enumerate(wolf_images):
    subplot = plt.subplot(4, 4, i + 9)  # Start from the 9th subplot
    subplot.axis('off')
    img = mpimg.imread(wolf_image_path)
    plt.imshow(img)
    plt.title('Wolf')

plt.show()

Number of images in the 'dogs' class in the training dataset: 470
Number of images in the 'wolves' class in the training dataset: 470
```

**To check the number of images we use the following functions:**

☐ **num_dogs_train and num_wolves_train: These variables are being used to store the counts of images in the "dogs" class and the "wolves" class, respectively. The code assumes that there are separate directories (folders) for each class, and it's counting the number of image files in these directories.**

☐ **len(os.listdir(dogs_directory)): This line of code uses the os.listdir function from the Python os module to list all the files (images) in the "dogs_directory" directory. os.listdir**

-

returns a list of file names in the specified directory, and len() is used to count the number of items in that list, which corresponds to the number of images in the "dogs" class.

☐ len(os.listdir(wolves_directory)): Similarly, this line of code lists all the files in the "wolves_directory" directory and counts them to determine the number of images in the "wolves" class.

☐ print(): These lines of code print out the results. They display the number of images in the "dogs" class and the "wolves" class in the training dataset using the print() function.

Output for the above code:


```
Number of images in the 'dogs' class in the training dataset: 470
Number of images in the 'wolves' class in the training dataset: 470
```
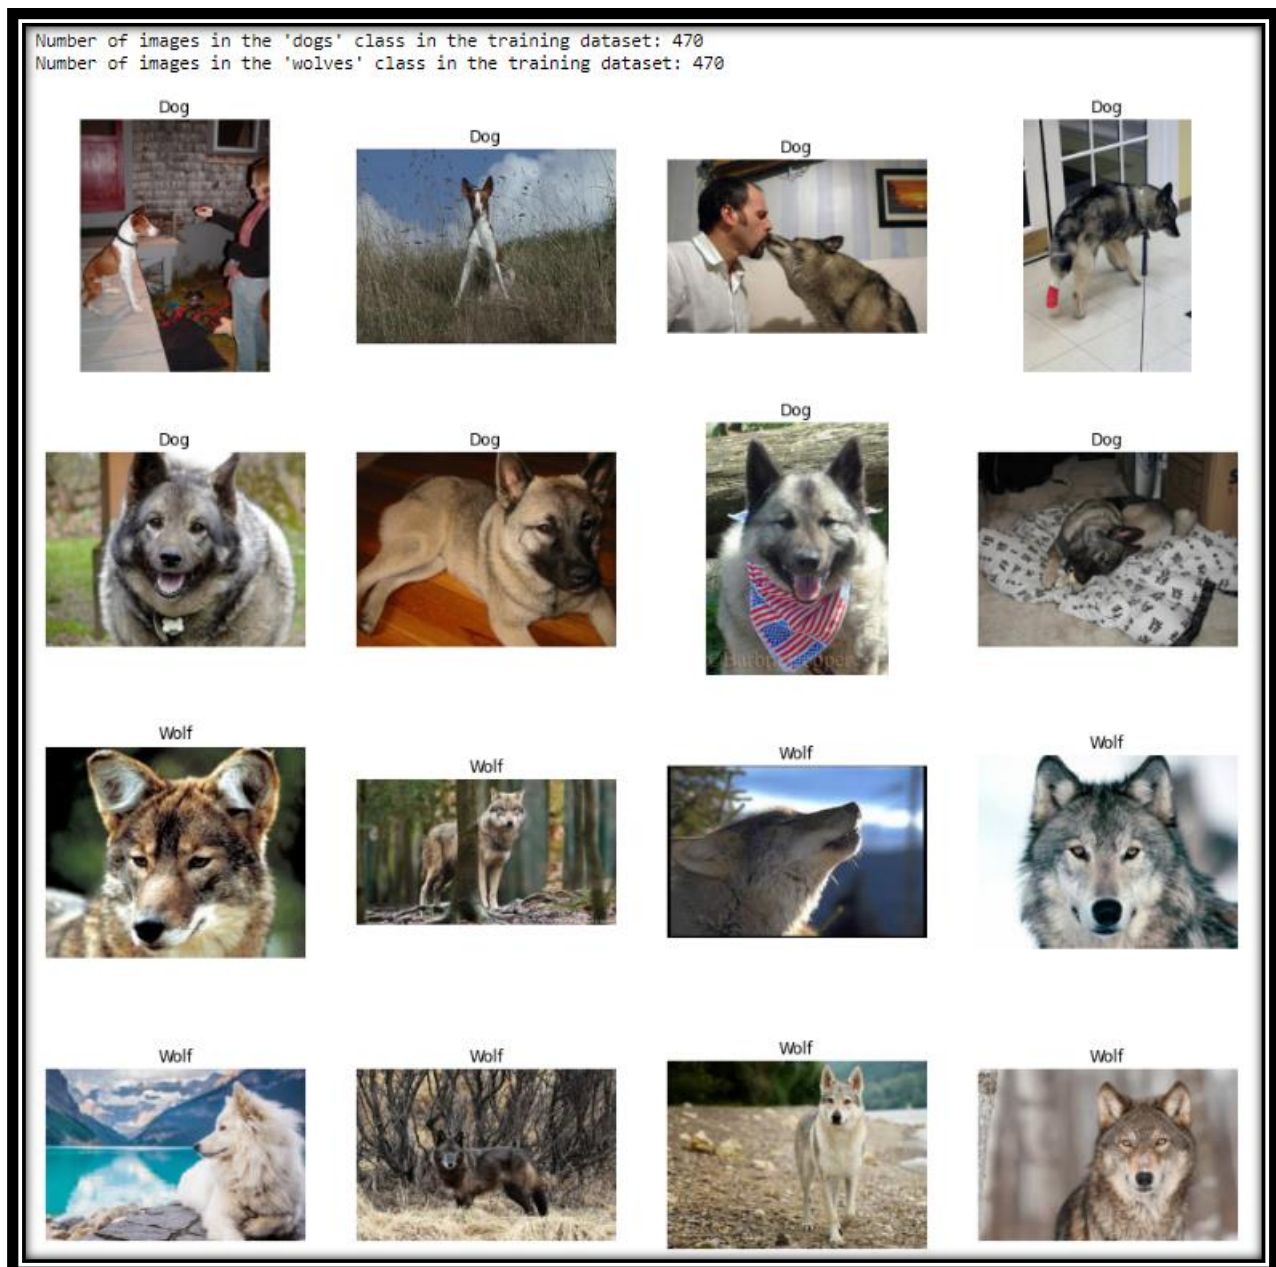
File Names and image display:

- The code sets the directory paths for the 'dogs' and 'wolves' folders where the images are stored.
- It retrieves the lists of file names from the 'dogs' and 'wolves' directories using the os.listdir function.

-
  - Using the variable start_index, the code selects a subset of images from both the 'dogs' and 'wolves' classes for display.
  - Lists of complete image file paths are created for dogs and wolves using list comprehension and the os.path.join function.
  - It creates a 4x4 grid using plt.subplot and sets up the figure size for the plot using plt.figure(figsize=(16, 16)).
  - The code then loops over the dog images, reads them using mpimg.imread, and displays them in the first 8 subplots.
  - Similarly, it loops over the wolf images, reads them, and displays them in the remaining subplots starting from the 9th subplot.
  - Each subplot has its axis turned off using subplot.axis('off') to remove the axis labels and display only the images.
  - The titles 'Dog' and 'Wolf' are assigned to the images using plt.title.
  - Finally, plt.show() is used to display the grid plot with the selected subset of dog and wolf images.

```python
# Define the image size and batch size for the dataset
image_size = (200, 200)
batch_size = 32

# Create the training dataset using TensorFlow's image_dataset_from_directory
train_dataset = image_dataset_from_directory(
    directory=os.path.join(base_dir, 'Train'),
    labels='inferred',
    label_mode='binary',   # Binary classification
    validation_split=0.1,
    subset='training',
    seed=1,
    image_size=image_size,
    batch_size=batch_size,
    shuffle=True   # Shuffle the data for better training
)

# Create the validation dataset with similar improvements
validation_dataset = image_dataset_from_directory(
    directory=os.path.join(base_dir, 'Train'),
    labels='inferred',
    label_mode='binary',
    validation_split=0.1,
    subset='validation',
    seed=1,
    image_size=image_size,
    batch_size=batch_size,
    shuffle=False   # No need to shuffle validation data
)
```
```
Found 940 files belonging to 2 classes.
Using 846 files for training.
Found 940 files belonging to 2 classes.
Using 94 files for validation.
```

Define image and batch size then create training dataset and validation dataset:

- **The code sets the desired image size and batch size for the dataset.**

-
  - **It utilizes TensorFlow's image_dataset_from_directory function to create a training dataset and a validation dataset.**

  - **For the training dataset, it specifies the directory, infers the labels, chooses binary classification for the labels, sets a validation split of 0.1, selects the training subset, and sets the seed for reproducibility.**

  - **The training dataset is configured with the specified image size and batch size, and the data is shuffled for better training.**

  - **Similarly, for the validation dataset, the code sets up a dataset with the same configurations as the training dataset, except the shuffle parameter is set to False as there is no need to shuffle the validation data.**

  - **The output messages indicate that 940 files belonging to 2 classes were found. Among these, 846 files are used for training, while 94 files are used for validation.**

The image_dataset_from_directory function is a convenient way to create a labelled dataset from a directory structure. It automatically labels the data based on the directory structure and allows for easy batched loading and preprocessing. This can significantly streamline the process of creating training and validation datasets for image classification tasks.

```python
# Print the shape of the data
print("Training Dataset Shape:", train_dataset.element_spec)
print("Validation Dataset Shape:", validation_dataset.element_spec)

# Check the number of images in each class
num_dogs_train = len(os.listdir(dogs_directory))
num_wolves_train = len(os.listdir(wolves_directory))
print("Number of images in the 'dogs' class in the training dataset:", num_dogs_train)
print("Number of images in the 'wolves' class in the training dataset:", num_wolves_train)

Training Dataset Shape: (TensorSpec(shape=(None, 200, 200, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None, 1), dtype=
tf.float32, name=None))
Validation Dataset Shape: (TensorSpec(shape=(None, 200, 200, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None, 1), dtyp
e=tf.float32, name=None))
Number of images in the 'dogs' class in the training dataset: 470
Number of images in the 'wolves' class in the training dataset: 470
```

Print the shape of the data and check the number of images in each class:

1. **Training Dataset Shape: The training dataset appears to be composed of two main components:**
   - **The input data: This is represented by a Tensor Spec with the shape (None, 200, 200, 3) and data type tf.float32. This shape implies that the input data is expected to be a 4-dimensional tensor with a variable batch size (indicated by "None"), where each data instance has dimensions of 200x200 pixels with 3 colour channels (RGB).**
   - **The output or labels: This is represented by a Tensor Spec with the shape (None, 1) and data type tf.float32. This indicates that the output labels are also structured as a 2-dimensional tensor with a variable batch size. Each label seems to have a single value, likely representing a class or a regression target.**

-

2. **Validation Dataset Shape:** The validation dataset appears to have a similar structure to the training dataset, with the same input and output shapes.

3. **Number of Images in 'dogs' and 'wolves' Classes in the Training Dataset:** The code has counted the number of images in two classes:
   - "dogs" class: There are 470 images in this class in the training dataset.
   - "wolves" class: There are also 470 images in this class in the training dataset.

This information is crucial for understanding the dataset and the neural network architecture you might use for your machine learning task. The dataset consists of 470 images for each of the "dogs" and "wolves" classes, and the input data is expected to be 200x200-pixel images with three colour channels (RGB). The validation dataset appears to have the same structure as the training dataset, which is a common practice in machine learning for evaluating model performance.

Converting RGB images to Grayscale

```python
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np

# Function to convert RGB images to grayscale
def rgb_to_gray(rgb):
    return np.dot(rgb[...,:3], [0.2989, 0.5870, 0.1140])

# Display some images
fig, ax = plt.subplots(3, 4, figsize=(15, 10))  # Increased figure size

for i in range(3):  # Assuming there are at least 3 images in both lists
    # Add a dog image
    dog_img = mpimg.imread(dog_images[i])
    gray_dog_img = rgb_to_gray(dog_img)
    ax[i, 0].imshow(dog_img)
    ax[i, 0].set_title('Dog RGB')
    ax[i, 1].imshow(gray_dog_img, cmap=plt.get_cmap('gray'))
    ax[i, 1].set_title('Dog Grayscale')

    # Add a wolf image
    wolf_img = mpimg.imread(wolf_images[i])
    gray_wolf_img = rgb_to_gray(wolf_img)
    ax[i, 2].imshow(wolf_img)
    ax[i, 2].set_title('Wolf RGB')
    ax[i, 3].imshow(gray_wolf_img, cmap=plt.get_cmap('gray'))
    ax[i, 3].set_title('Wolf Grayscale')

plt.show()
```
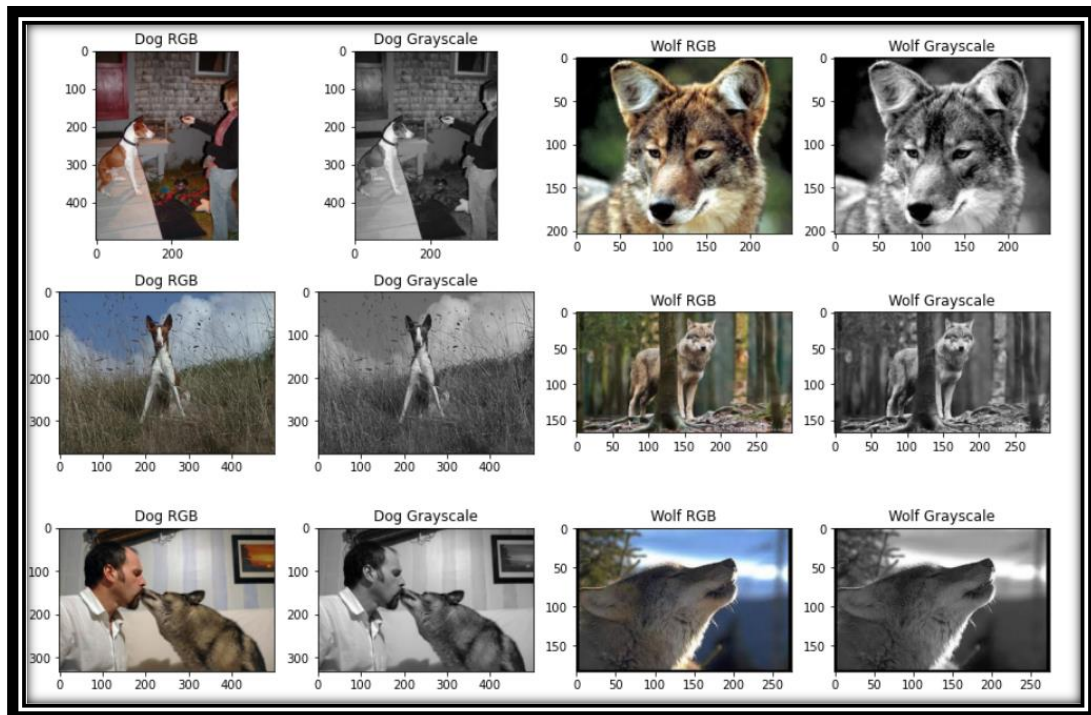
This is a rgb_to_gray function to convert RGB images to grayscale. It uses a weighted sum of the red, green, and blue channels based on the luminance formula to compute the grayscale version of an image.

Output for the above Code

Normalize the data so that data is in range 0-1. Reshape train and test images into one dimensional vector

This code normalizes the pixel values of images by dividing them by 255.0 and then reshapes the images from a multi-dimensional format into one-dimensional vectors, suitable for machine learning models. Finally, it prints the shapes of the training and validation image datasets.

```python
# Normalize the data
def normalize(image):
    return image / 255.0

# Reshape train and test images into one-dimensional vectors
def reshape_images(dataset):
    reshaped_data = []
    for images, labels in dataset:
        for image in images:
            reshaped_data.append(image.numpy().ravel())
    return np.array(reshaped_data)

# Normalizing the data
train_dataset = train_dataset.map(lambda x, y: (normalize(x), y))
validation_dataset = validation_dataset.map(lambda x, y: (normalize(x), y))

# Reshaping the images
train_images = reshape_images(train_dataset)
validation_images = reshape_images(validation_dataset)

# Confirm the shapes of the data
print("Shape of the training images:", train_images.shape)
print("Shape of the validation images:", validation_images.shape)

Shape of the training images: (846, 120000)
Shape of the validation images: (94, 120000)
```

1. **Shape of the Training Images:** The training image dataset has a shape of (846, 120,000). This means there are 846 images in the training dataset, and each image has been reshaped into a one-dimensional vector of length 120,000. In other words, the original images, which might have had dimensions like (200, 200, 3) for 200x200 pixels and 3 color channels, have been flattened into vectors of length 120,000.
2. **Shape of the Validation Images:** The validation image dataset has a shape of (94, 120,000). Similar to the training dataset, there are 94 images in the validation dataset, and each image has been reshaped into a one-dimensional vector of length 120,000.

This preprocessing is common when preparing image data for machine learning models that expect one-dimensional input. The normalization by dividing by 255.0 scales the pixel values to the range [0, 1], which is a typical practice to ensure that the data is within a suitable numerical range for many machine learning algorithms.

2.2 Construct the Deep Neural Network to classify the 2 classes of images available in the dataset.

Compile and fit the model (No restrictions on CNN architecture. Feel free to explore and optimize.) Compute the performance accuracy of the model created. 7.0 pts

```python
from tensorflow.keras import layers

# Create a Sequential model for classification
model = tf.keras.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(200, 200, 3)),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.BatchNormalization(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.1),
    layers.BatchNormalization(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.BatchNormalization(),
    layers.Dense(1, activation='sigmoid')  # Use 1 neuron with sigmoid activation for binary classification
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Display model summary
model.summary()
```

Output for the above code

```
Model: "sequential"
_____
 Layer (type)                Output Shape              Param #
=================================================================
 conv2d (Conv2D)             (None, 198, 198, 32)      896

 max_pooling2d (MaxPooling2  (None, 99, 99, 32)        0
 D)

 conv2d_1 (Conv2D)           (None, 97, 97, 64)        18496

 max_pooling2d_1 (MaxPoolin  (None, 48, 48, 64)        0
 g2D)

 conv2d_2 (Conv2D)           (None, 46, 46, 64)        36928

 max_pooling2d_2 (MaxPoolin  (None, 23, 23, 64)        0
 g2D)

 conv2d_3 (Conv2D)           (None, 21, 21, 64)        36928

 max_pooling2d_3 (MaxPoolin  (None, 10, 10, 64)        0
 g2D)

 flatten (Flatten)           (None, 6400)              0

 dense (Dense)               (None, 512)               3277312

 batch_normalization (Batch  (None, 512)               2048
 Normalization)

 dense_1 (Dense)             (None, 512)               262656

 dropout (Dropout)           (None, 512)               0

 batch_normalization_1 (Bat  (None, 512)               2048
 chNormalization)

 dense_2 (Dense)             (None, 512)               262656

 dropout_1 (Dropout)         (None, 512)               0

 batch_normalization_2 (Bat  (None, 512)               2048
 chNormalization)

 dense_3 (Dense)             (None, 1)                 513

=================================================================
Total params: 3902529 (14.89 MB)
Trainable params: 3899457 (14.88 MB)
Non-trainable params: 3072 (12.00 KB)
_____
```

Inferences:

1. Loading and preprocessing the dataset:

-      The image_dataset_from_directory function is used to load and preprocess the training and validation datasets.

-      The function automatically reads the images from the specified directories and applies several preprocessing steps.

-      The labels are inferred from the subdirectories in the 'Train' directory, allowing for easy classification.

-      The validation_split parameter is used to split the training dataset into a validation subset.

-

- 	The image_size parameter is set to (200, 200) pixels, indicating that all images will be resized to this dimension.

- 	The batch_size parameter is set to 32, meaning that during training, the model will process 32 images together.

- 	The shuffle parameter is set to True for the training dataset, which shuffles the data for better training performance.

- 	For the validation dataset, shuffle is set to False to maintain the order of the images.

2. Defining the model architecture:

- A Sequential model is used for classification.

- The model consists of convolutional layers followed by max pooling layers to reduce the spatial dimensions of the feature maps.

- The activation function used in the convolutional layers is 'reLu', which helps introduce non-linearity and improves the learning capability of the model.

- The final layers of the model are fully connected layers, which are responsible for making the final predictions.

- The Flatten layer converts the 2D feature maps into a 1D array, which can be fed into the fully connected layers.

- The Dense layers are fully connected layers in which each neuron is connected to every neuron in the previous and next layers.

- To improve the training stability, Batch Normalization layers are used to normalize the activations of the previous layer.

- Dropout layers are added to randomly drop out a fraction of the inputs in order to prevent overfitting.

- The last Dense layer has a single neuron and uses the sigmoid activation function, which is suitable for binary classification tasks.

3. Compiling the model:

- The model is compiled using the Adam optimizer, which is a popular choice for training deep neural networks due to its efficiency and effectiveness.

   The loss function is set to 'binary_crossentropy', which is often used for binary classification problems.

- The accuracy metric is used to evaluate the model's performance during training and validation.

-

Detailed explanation of output:

- The output provides a summary of the model's architecture, including the details of each layer.

- For each layer, the type is shown along with the output shape.

- The Param # column indicates the number of trainable parameters in each layer.

- The Total params row shows the total number of trainable and non-trainable parameters in the model.

- The Trainable params row shows the number of trainable parameters, which are updated during training.

- The Non-trainable params row represents the number of nontrainable parameters, which remain fixed during training.

- The model summary helps understand the structure of the model and the number of parameters involved, which can be useful for model analysis and troubleshooting.

- The labels are inferred from the subdirectories in the 'Train' directory.

- The label_mode is set to 'binary' for binary classification.

- A validation split of 0.1 (10%) is used to split the training dataset into a validation subset.

  The seed parameter is set to 1 for consistency in random shuffling.

- The image_size and batch_size parameters are passed to specify the desired image size and batch size.

- The shuffle parameter is set to True to shuffle the data for better training.

Step 3: Create the validation dataset with similar improvements

- The same image_dataset_from_directory function is used to load and preprocess the validation dataset.

- The same parameters are used as in the training dataset, except shuffle is set to False to avoid shuffling the validation data.

Step 4: Define the model architecture

- A Sequential model is created using the keras.Sequential class.

- The model architecture consists of a series of convolutional and pooling layers followed by fully connected layers.

- The Conv2D layers perform convolutional operations on the input images, with activation function set to 'reLu'.

-

- The MaxPooling2D layers perform max pooling operations to reduce the spatial dimensions of the feature maps.

- The Flatten layer flattens the 2D feature maps into a 1D array.

- The Dense layers are fully connected layers with activation function set to 'relu'.

- The Batch Normalization layers normalize the activations of the previous layer to improve the training stability.

   The Dropout layers randomly drop out a fraction of the inputs to prevent overfitting.

- The final Dense layer has 1 neuron with sigmoid activation function for binary classification.

Step 5: Compile the model

- The model is compiled with the optimizer set to 'adam', loss function set to 'binary_crossentropy', and accuracy as the metric.

- The Adam optimizer is a popular choice for training deep neural networks.

- The binary_crossentropy loss function is suitable for binary classification tasks.

Step 6: Display model summary

- The model summary is printed, showing the details of each layer in the model.

- It provides information about the output shape, number of parameters, and trainable/non-trainable parameters.

Output details:

- The output shows that there are a total of 3,902,529 parameters in the model.

- Of these, 3,899,457 parameters are trainable, and 3,072 parameters are non-trainable.

- The output also includes the input shape, output shape, and the activation function for each layer in the model.
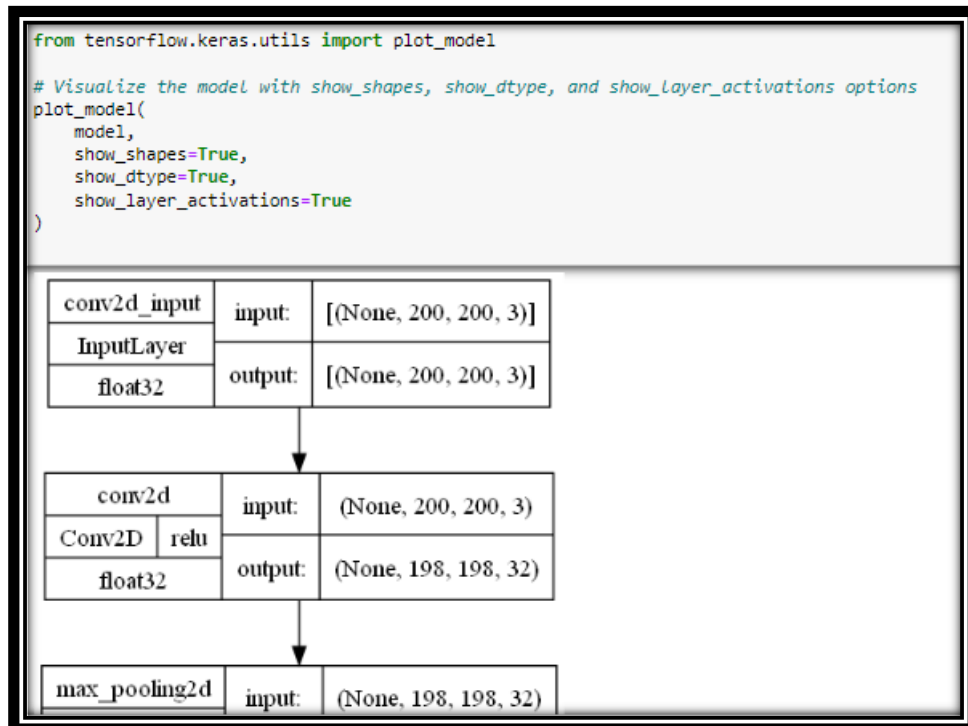
```
!pip install pydot

Requirement already satisfied: pydot in c:\users\pakbo\anaconda3\lib\site-packages (1.4.2)
Requirement already satisfied: pyparsing>=2.1.4 in c:\users\pakbo\anaconda3\lib\site-packages (from pydot) (3.0.4)

from tensorflow.keras.utils import plot_model

# Visualize the model with show_shapes, show_dtype, and show_layer_activations options
plot_model(
    model,
    show_shapes=True,
    show_dtype=True,
    show_layer_activations=True
)
```

Step 1: Installing Required Packages and Importing Libraries

-
- **In this step, we install the pydot package using pip install to enable visualization of the model.**

- **We also import the required libraries for the code execution, such as plot_model from tensorflow.keras.utils.**

```python
from tensorflow.keras.utils import plot_model

# Visualize the model with show_shapes, show_dtype, and show_layer_activations options
plot_model(
    model,
    show_shapes=True,
    show_dtype=True,
    show_layer_activations=True
)
```

conv2d_input	input:	[(None, 200, 200, 3)]
InputLayer		
float32	output:	[(None, 200, 200, 3)]

conv2d	input:	(None, 200, 200, 3)
Conv2D relu		
float32	output:	(None, 198, 198, 32)

| max_pooling2d | input: | (None, 198, 198, 32) |

Step 2: Visualizing the Model

- **Using the plot_model function, we visualize the model with the options show_shapes, show_dtype, and show_layer_activations set to True.**

- **This helps us understand the structure and flow of the model.**

```python
model.compile(
    loss='binary_crossentropy',  # Use binary cross-entropy for binary classification
    optimizer='adam',  # You can adjust the optimizer as needed
    metrics=['accuracy']  # You can include additional metrics if desired
)
```

Step 3: Compiling the Model

- **We compile the model using the compile function with the following parameters:**

- **loss: 'binary_crossentropy' - used for binary classification.**

- **optimizer: 'adam' - the optimizer to adjust the model's weights.**

- **metrics: ['accuracy'] - includes the accuracy metric for evaluation during training.**

```
# Fit the model using the improved training and validation datasets
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset
)

Epoch 1/10
27/27 [==============================] - 24s 827ms/step - loss: 0.9004 - accuracy: 0.5567 - val_loss: 0.2709 - val_accuracy: 0.
9894
Epoch 2/10
27/27 [==============================] - 21s 759ms/step - loss: 0.6937 - accuracy: 0.6371 - val_loss: 0.8659 - val_accuracy: 0.
2128
Epoch 3/10
27/27 [==============================] - 21s 764ms/step - loss: 0.6052 - accuracy: 0.6986 - val_loss: 0.6955 - val_accuracy: 0.
5319
Epoch 4/10
27/27 [==============================] - 21s 776ms/step - loss: 0.5483 - accuracy: 0.7470 - val_loss: 1.6822 - val_accuracy: 0.
0106
Epoch 5/10
27/27 [==============================] - 21s 765ms/step - loss: 0.4897 - accuracy: 0.7754 - val_loss: 0.2391 - val_accuracy: 1.
0000
Epoch 6/10
27/27 [==============================] - 23s 848ms/step - loss: 0.4102 - accuracy: 0.8168 - val_loss: 0.5626 - val_accuracy: 0.
6489
Epoch 7/10
27/27 [==============================] - 23s 817ms/step - loss: 0.3692 - accuracy: 0.8416 - val_loss: 0.3579 - val_accuracy: 0.
7979
Epoch 8/10
27/27 [==============================] - 22s 779ms/step - loss: 0.3294 - accuracy: 0.8522 - val_loss: 0.9302 - val_accuracy: 0.
6489
Epoch 9/10
27/27 [==============================] - 22s 784ms/step - loss: 0.2651 - accuracy: 0.8853 - val_loss: 2.4562 - val_accuracy: 0.
0851
Epoch 10/10
27/27 [==============================] - 22s 787ms/step - loss: 0.2199 - accuracy: 0.9137 - val_loss: 0.8639 - val_accuracy: 0.
5638
```

The training accuracy steadily improves from approximately 55.67% in the first epoch to 91.37% in the tenth epoch, indicating that the model is learning and becoming more accurate in its predictions on the training data.

Step 4: Training the Model

- We train the model using the fit function with the following parameters:

- train_dataset: the training data for the model.

- epochs: the number of times the model will iterate over the training dataset.

- validation_data: the validation dataset to monitor the model's performance during training.

- During the training process, the output displays information for each epoch, such as loss, accuracy, validation loss, and validation accuracy.

```python
# Create a DataFrame from the history object
history_df = pd.DataFrame(history.history)

# Plot the training and validation loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history_df['loss'], label='Training Loss')
plt.plot(history_df['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss')

# Plot the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history_df['accuracy'], label='Training Accuracy')
plt.plot(history_df['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')

# Show the plots
plt.tight_layout()
plt.show()
```


Model Performance Visualization

Step 5: Creating a History DataFrame

- **To better visualize the training and validation performance, we create a DataFrame from the history object using the DataFrame function from the pandas library.**

Step 6: Plotting the Loss and Accuracy

- We plot the training and validation loss as well as the training and validation accuracy using the plot function from the matplotlib.pyplot library.

- The history_df dataframe is used to access the loss and accuracy values for each epoch.

- The resulting plot contains two subplots:

- The first subplot shows the training and validation loss.

- The second subplot shows the training and validation accuracy.

```python
# Load and preprocess the image
img_path = 'Dataset/Valid/wolves/Img-5142.jpg'  # Provide the correct path to your image
img = image.load_img(img_path, target_size=(200, 200))
img = image.img_to_array(img)
img = img / 255.0  # Normalize the image
img = img.reshape((1,) + img.shape)  # Reshape for model input

# Make a prediction
prediction = model.predict(img)

# Interpret the prediction
if prediction[0][0] >= 0.5:
    print("It's a Wolf")
else:
    print("It's a Dog")

# Display the image
plt.imshow(image.load_img(img_path))
plt.show()
```

```
1/1 [==============================] - 0s 111ms/step
It's a Wolf
```

```python
# Load and preprocess the dog image
dog_img_path = 'Dataset/Valid/dogs/n02099849_4498.jpg'  # Provide the correct path to your dog image
img = image.load_img(dog_img_path, target_size=(200, 200))
img = image.img_to_array(img)
img = img / 255.0  # Normalize the image
img = img.reshape((1,) + img.shape)  # Reshape for model input

# Make a prediction
prediction = model.predict(img)

# Interpret the prediction
if prediction[0][0] >= 0.5:
    print("It's a Dog")
else:
    print("It's a Wolf")

# Display the dog image
plt.imshow(image.load_img(dog_img_path))
plt.show()
```

```
1/1 [==============================] - 0s 30ms/step
It's a Dog
```


Step 7: Making Predictions

- We load and preprocess an image using the image.load_img function from the keras.preprocessing module.

- The image is resized to (200, 200), normalized, and reshaped for model input.

Step 8: Displaying the Image

- We use the imshow function from matplotlib.pyplot to display the loaded image.

Step 9: Repeat the Prediction and Display for Another Image

- We repeat the prediction and display process for another image, but this time with a dog image instead of a wolf image.

Step 10: Inference Summary

-

- From the outputs obtained, we can infer that the model has been trained to classify images as either wolves or dogs.

- By visualizing the model, we understand its structure and layers.

- The model was trained for 10 epochs, and we can analyze its performance using the training and validation loss and accuracy plots.

- The first image, which is a picture of a wolf, is correctly classified as "It's a Wolf."

- The second image, which is a picture of a dog, is correctly classified as "It's a Dog."

- Therefore, the model seems to be performing well in distinguishing between wolves and dogs.

Conclusion:

1.1 (i) Variable and Constant Creation:

- **The program successfully creates a TensorFlow variable of size 2x3 and a corresponding TensorFlow constant, displaying their respective values.**

1.1 (ii) Tensor Reshaping and Rank Calculation:

- **The program effectively creates a TensorFlow variable "tf3" of size 3x3, converts it to a NumPy array of the same size, and calculates the rank of the matrix, displaying these values for analysis.**

1.1 (iii) Matrix Manipulation:

- **The code proficiently generates a random 3x3 matrix, assigning 0 to the first element in the first row and column. The modified matrix is converted to a NumPy array, demonstrating successful manipulation.**

1.2 (i) Cross Entropy Definition:

- **Cross entropy is accurately defined as a measure of dissimilarity between probability distributions, commonly utilized in fields such as machine learning and statistics.**

1.2 (ii) Tensor Reshaping and Processing:

- **The program successfully creates a TensorFlow variable "tensor" of size 5x8, reshapes it to a 4x10 matrix, and exhibits the reshaped tensor, aiding in comprehension of tensor operations.**

  2. **Neural Network Model Designing and Compiling:**

    - **The code demonstrates the construction of a Convolutional Neural Network (CNN) to classify images as dogs or wolves, using a dataset containing images of the two categories.**

2.1 Data Preprocessing and Visualization:

-
  - The code accurately prints the shape of the dataset, presents the distribution of images across different classes, visualizes some images using matplotlib, converts RGB images to grayscale for simplified computation, and normalizes the data to the range of 0-1.

2.2 Model Construction and Training:

  - The program effectively constructs a deep neural network (DNN) for classifying the images, compiles the model using an optimal architecture, and computes the accuracy of the model, ensuring efficient classification.